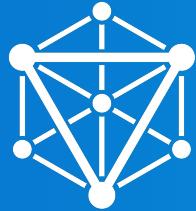




VISIUM

Cutting-edge AI for the world's best brands

Spinoff ETHzürich



Shedding light on graph neural networks

AMLD Workshop - 29.09.2021

Arnaud Dhaene & William Cappelletti

Your Visium team for today



William Cappelletti
Machine Learning Engineer



Arnaud Dhaene
AI Solution Specialist



Camilla Casamento
Machine Learning Engineer



Matteo Togninalli, PhD
Chief Operating Officer

Visium | Leading AI-consultancy with close connection to academia

Broad experience in AI/ML projects

Successfully completed 90+ projects with 40+ happy clients of combined revenues over CHF 400 BN (Healthcare: 20+ projects).



One-stop-shop for AI

Expertise in providing E2E AI-service from use-case ideation, model development to scaling within organization.

Visium's competitive advantage

Exponential growth

Advising leading companies on AI since 2018 and grown to 40 team members without external funding.

Academia collaboration

Visium is an ETH spin-off and retains close collaboration with leading AI-universities.



ETH zürich

EPFL

Stanford University

HARVARD
UNIVERSITY

Imperial College
London

VISIUM

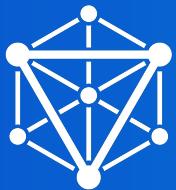


Table of Contents

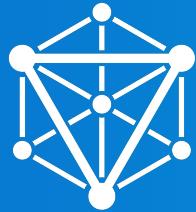
- Introduction
- Graph data
- Deep learning on Graphs
- Interpretability

>_

>_

>_

VISIUM



Introduction

Deep Learning on Graphs?

Introduction to Deep Learning

Machine Learning studies algorithms
that improve by **observing data**

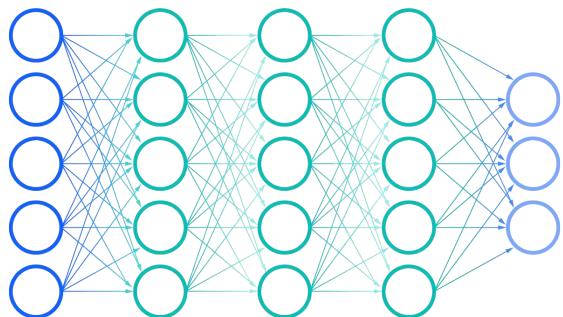
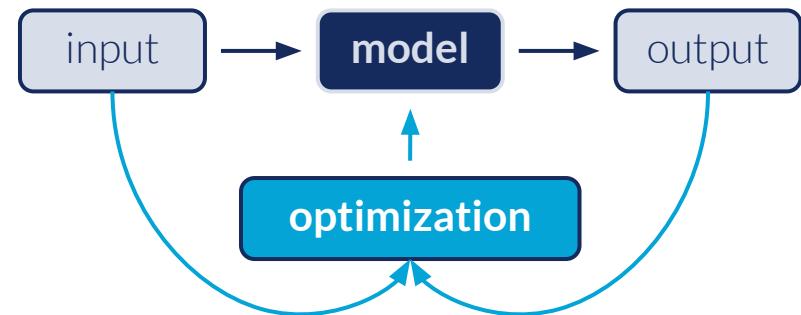


Fig. from ibm.com/cloud/learn/neural-networks

Deep Learning is a family of ML methods based on
Artificial Neural Networks.

It has been very successful thanks to its **flexibility** to
adapt **architectures** and various **data types** to the
specific **task**.

Data contains structure that can be leveraged

Data have inherent **structure**, which should be taken into account by the models:

- Images have spatial relationships
→ CNNs (convolutions)

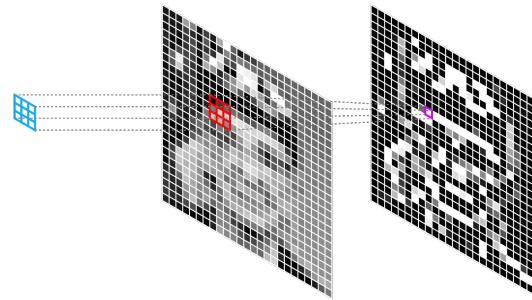


Fig. from [Gregory Gundersen blog](#)

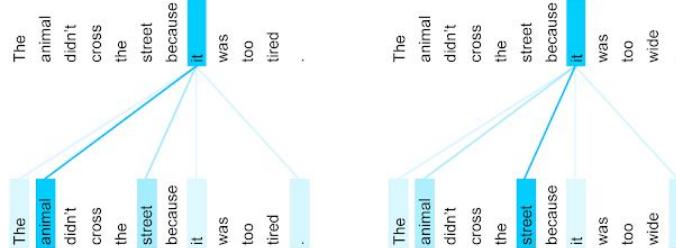


Fig. from [Google AI blog](#)

- Language: sequential and syntactic structures
→ Transformers (attention and recurrence)

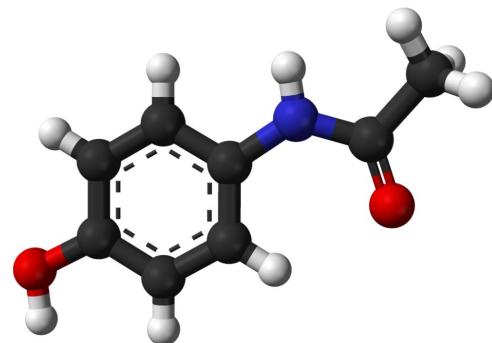
What if our data has even more structure?

Network representations contain a wealth of information

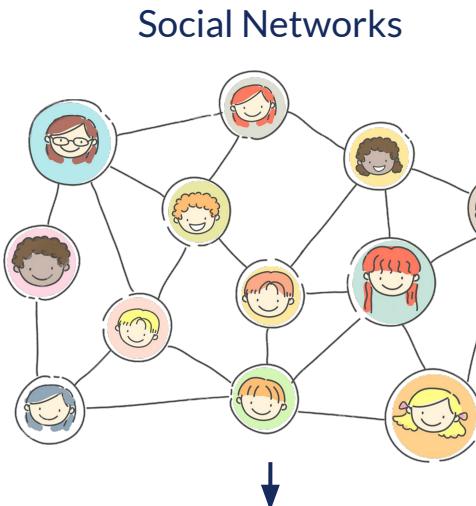
Many real world data can be naturally modeled as **graphs**:

Sets of entities, or **nodes**, with pairwise relations, i.e. a set of **edges**

Molecules



↓
Atoms sharing bonds



Transportation Maps

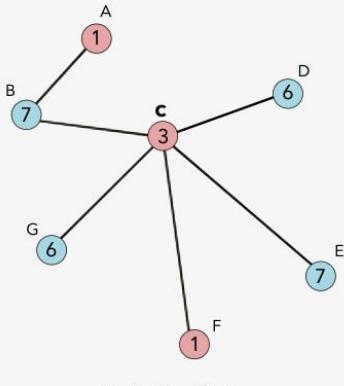


Intersections connected by roads

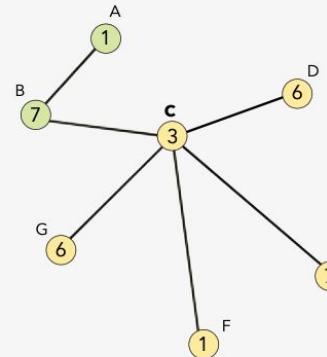
VISIUM

Network representations are used to solve problems

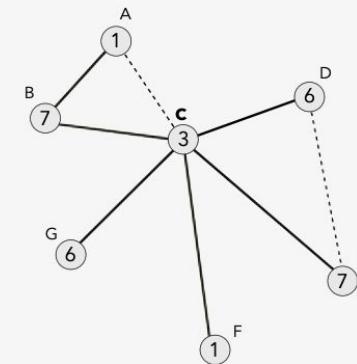
There are **many useful problems** that can be formulated over graphs. The structure's flexibility allows for a wide problem setting. The most common ones are:



Toxic



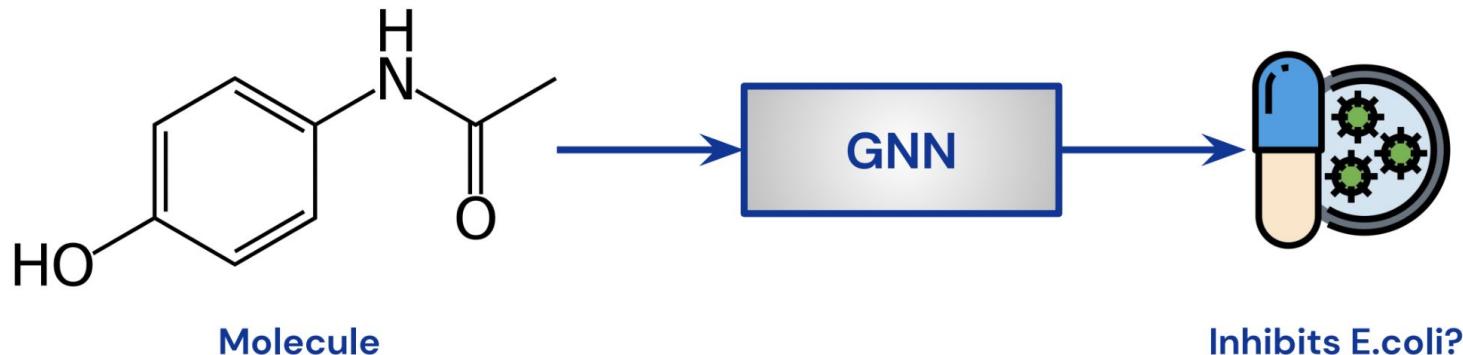
Node Clustering



Graph Neural Networks are the method of choice for Graphs

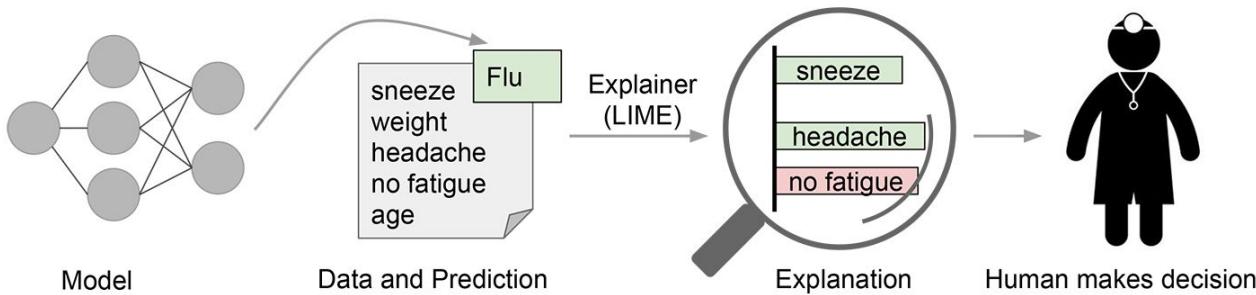
Previous Deep Learning architectures (CNNs or Transformers) DO NOT work on graphs

- We have to introduce specific **Graph Neural Networks (GNN)**



GNNs, like other Neural Networks, are *Black Box Models*

More complex models make it harder to see what goes on under the hood.
This holds for GNNs as well.



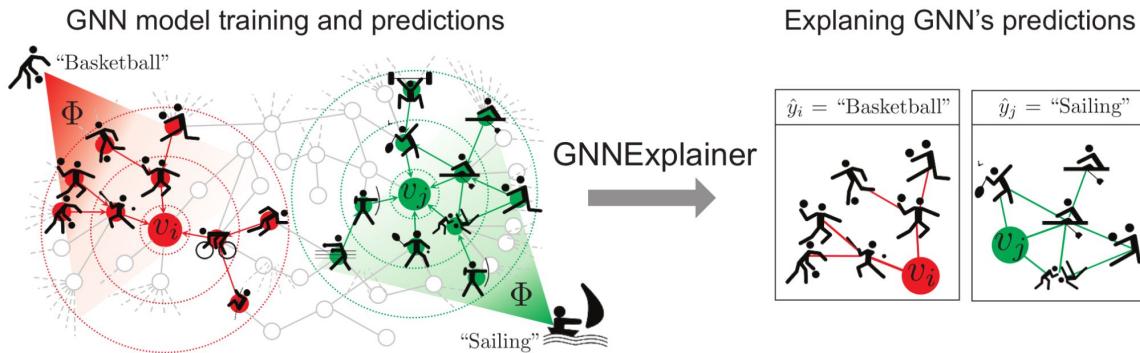
In many applications, for instance *medical* or *decisional* problems, the **motivation** of an answer is as important as the answer itself, if not more.

We need **accountability** and **transparency**.

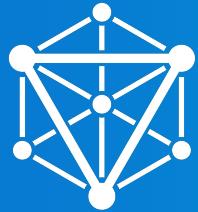
Explaining the outcome of GNNs is essential to their use

We should try to **explain the predictions** made by the complex GNN models:

- **Why** does the GNN predict a certain output?
- **How** do features influence the prediction?
- **What** structures are semantically relevant for a certain *sample*, or a given class?



→ Part of these questions are answered by **GNNExplainer**

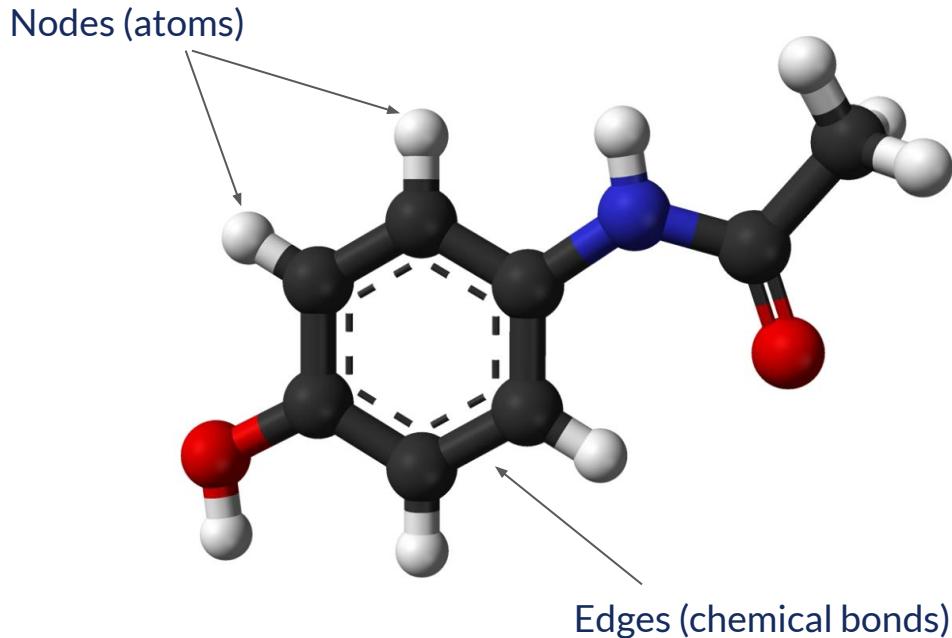


Graph data

Instilling natural structure in data

Network representations contain a wealth of information

A graph $G = (V, E)$ can carry information both in its **nodes** and **edges**

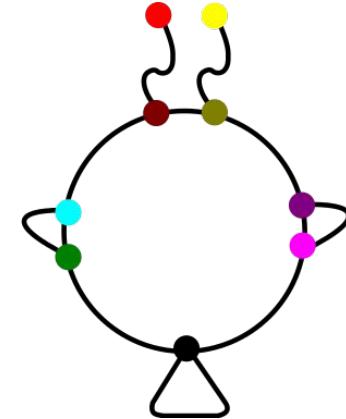
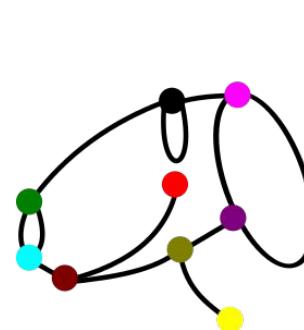
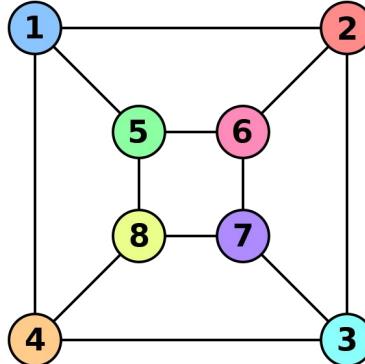
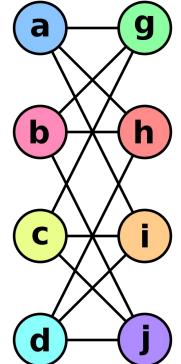


- Node features
e.g. atomic number
- Edge features
e.g. type of bond
(simple, double)

Graphs have very specific properties

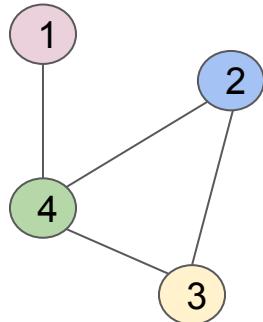
The main property of graphs is in **how the nodes are connected**

Graphs **structure** is **independent** from the labelling of the nodes, or from how we choose to draw them

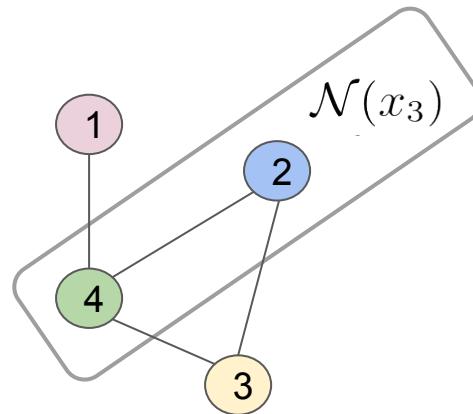


Graph representations

Adjacency matrix: $A_{i,j} = 1$ iff there is an edge from node i to node j , otherwise $A_{i,j} = 0$. Traditional network analysis techniques heavily rely on the use of the adjacency matrix.



$$A = \begin{bmatrix} 0 & 0 & 0 & 1 \\ 0 & 0 & 1 & 1 \\ 0 & 1 & 0 & 1 \\ 1 & 1 & 1 & 0 \end{bmatrix}$$



Neighbourhood: $\mathcal{N}(x_v)$ represents the set of neighbors of node v

Edge list: a list E of edges, each represented as a pair (i, j) , meaning that node i is connected to node j .

$$E = [[1, 4], [4, 1], [2, 4], [4, 2], [2, 3], [3, 2], [3, 4], [4, 3]]$$

There are sophisticated libraries to work with graph data

There are many libraries to integrate Graphs in ML frameworks



We will use PyTorch Geometric

- + Adjacency and features matrices as (sparse) Tensors
- + Includes many datasets
- + Supports batching and data loaders

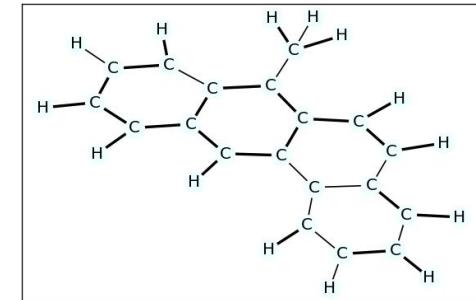
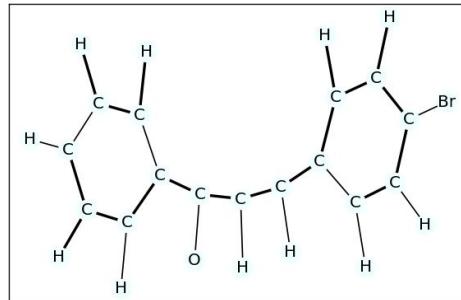
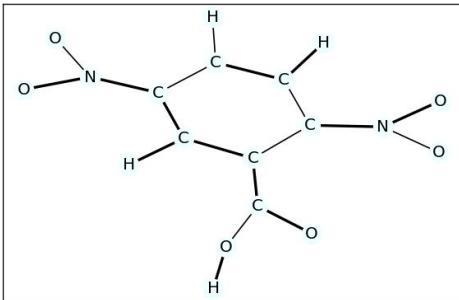


Let's get our hands dirty with some graphs!



Colab Notebook

1. Data loading and EDA
2. Graph visualization
3. Batching



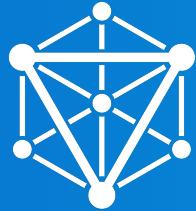
Open the notebook on Colab



1. Go to <https://colab.research.google.com/>
2. Under GitHub search for VisiumCH
3. Choose the AMLD-2021-Graphs repo
4. Open notebooks/workshop_notebook.ipynb

A screenshot of the Google Colab interface. At the top, there's a navigation bar with tabs: Examples, Recent, Google Drive, GitHub (which is currently selected), and Upload. Below the navigation bar, there's a search bar with the placeholder "Enter a GitHub URL or search by organization or user". To the right of the search bar is a checked checkbox labeled "Include private repos". The search term "VisiumCH" is entered into the search bar. Below the search bar, there are two dropdown menus: "Repository:" set to "VisiumCH/AMLD-2021-Graphs" and "Branch:" set to "master". There are also "Path" and "Path" dropdowns. The main content area shows two items listed: "notebooks/workshop_notebook.ipynb" and "notebooks/workshop_notebook_no_sol.ipynb", each with a GitHub icon, a magnifying glass icon for search, and a copy icon.

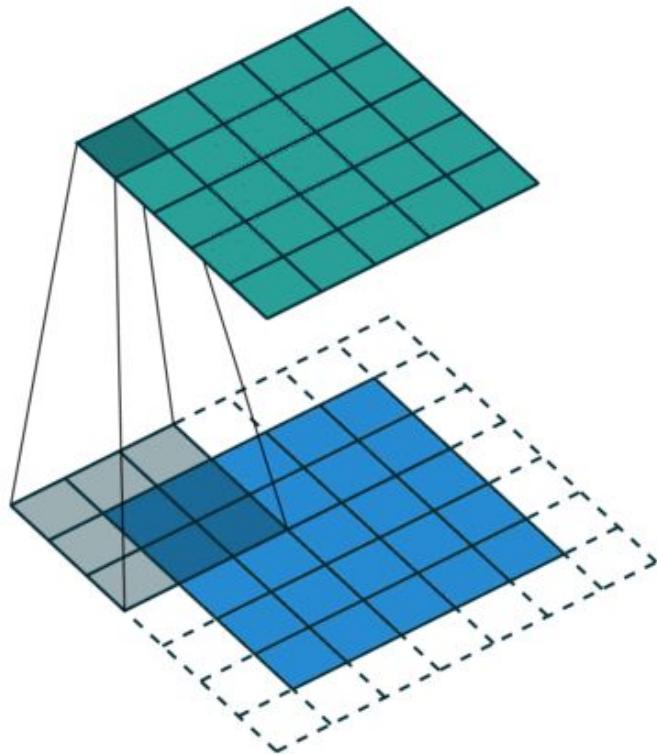
VISIUM



Deep Learning on Graphs

The implementation of Graph Neural Networks

Leveraging spatial information with convolutions



Properties of image convolution

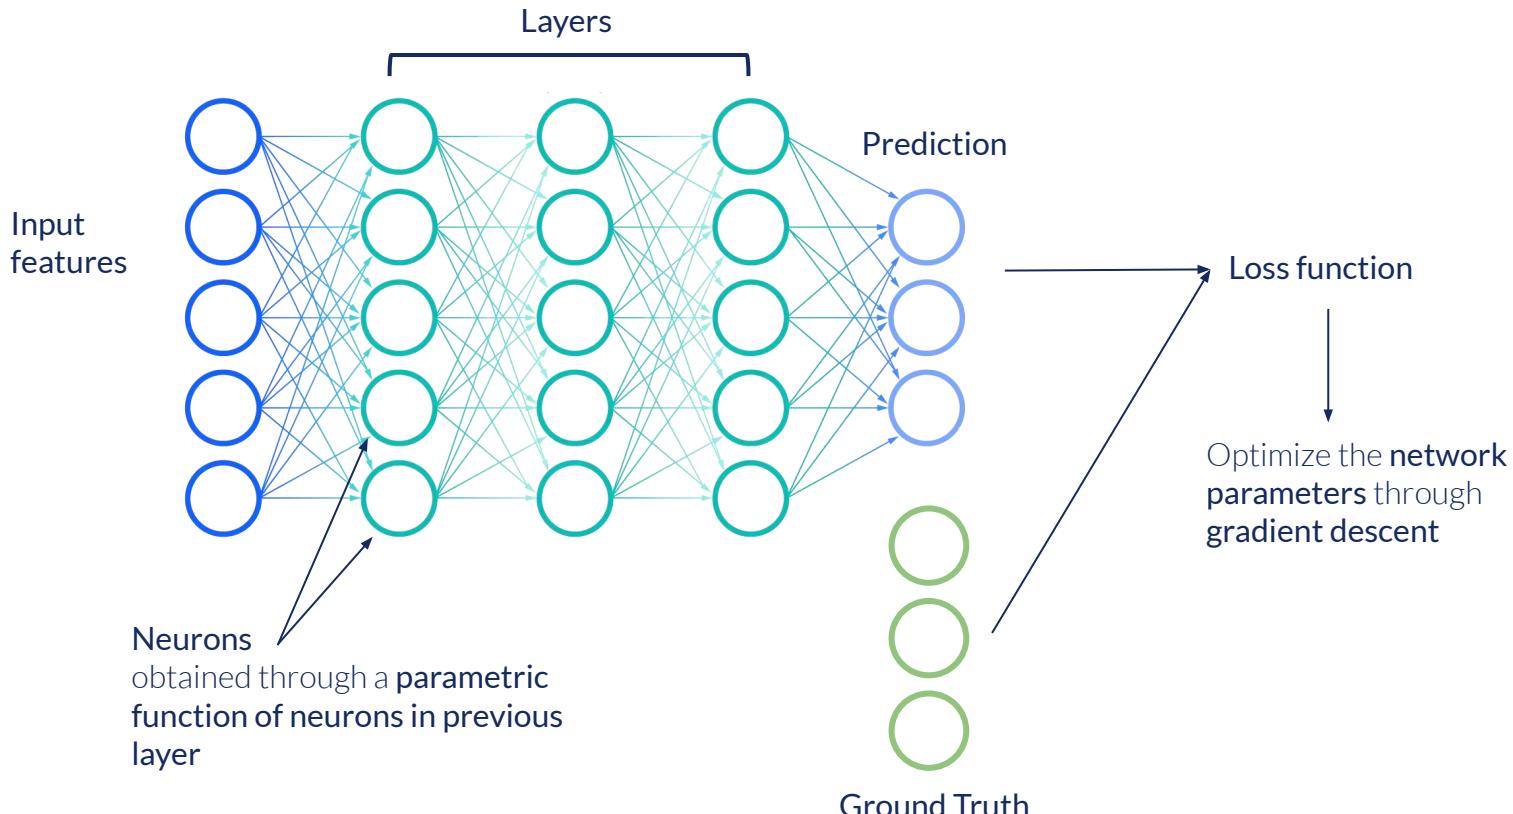
- **Locality**

Exploit position information

- **Translation equivariance**

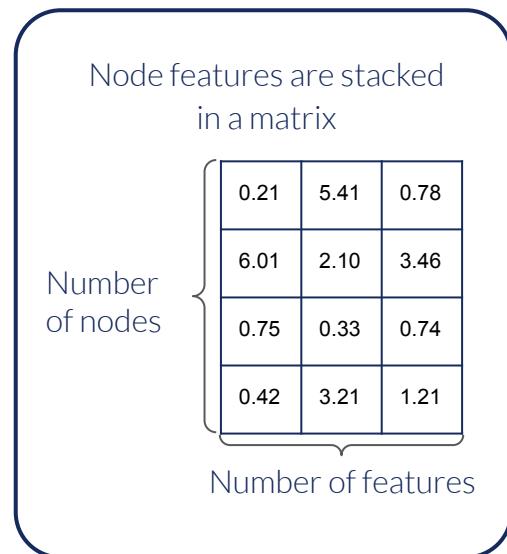
A translation of the input is reflected in the output (e.g. image segmentation)

The general structure of artificial neural networks



Towards Graph Neural Networks

What are the desired properties that graph convolution should satisfy?



This introduces an ordering of the nodes

Graph classification

The GNN output should not depend on such ordering

Node classification

The GNN output should not depend on the ordering but reflect it exactly on output

Permutation invariance and equivariance

Let X be our feature matrix, P a permutation matrix and A the graph adjacency matrix.

- A function f is said to be **permutation invariant** if

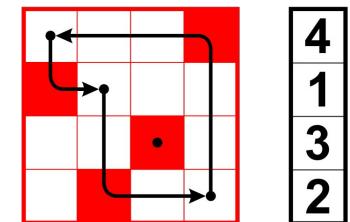
$$f(PX, PAP^T) = f(X, A)$$

Schema of a permutation P
acting on a column vector

1
2
3
4

- A function f is said to be **permutation equivariant** if

$$f(PX, PAP^T) = P(f(X, A))$$



The desired properties for Graph Neural Networks

- **Locality:** the features of the neighbors of the considered node i , $N(x_i)$, should influence the computation of the latent representation of i
- **Permutation equivariance:** if we permute the order of the nodes in the input matrix, such permutation will also be applied to the nodes latent representations

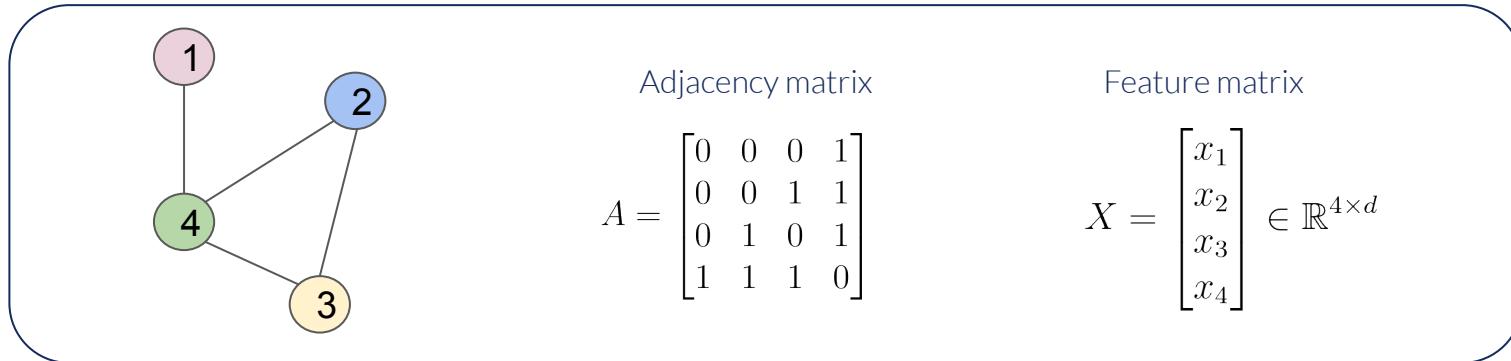
$$f(X, A) = \begin{pmatrix} g(x_1, \mathcal{N}(x_1)) \\ g(x_2, \mathcal{N}(x_2)) \\ \vdots \\ g(x_N, \mathcal{N}(x_N)) \end{pmatrix}$$

graph convolution function f ,
permutation equivariant

g is permutation invariant
over $N(x_i)$

the ordering of the
neighbors of a node
should not influence the
node's latent
representation

The inner workings of Graph Neural Networks | Example

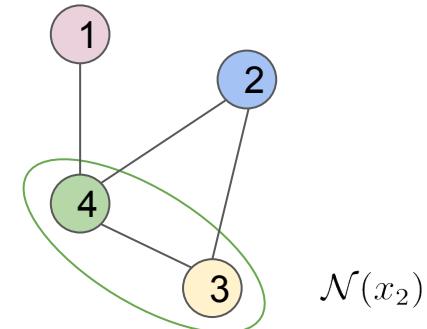


Graph convolution

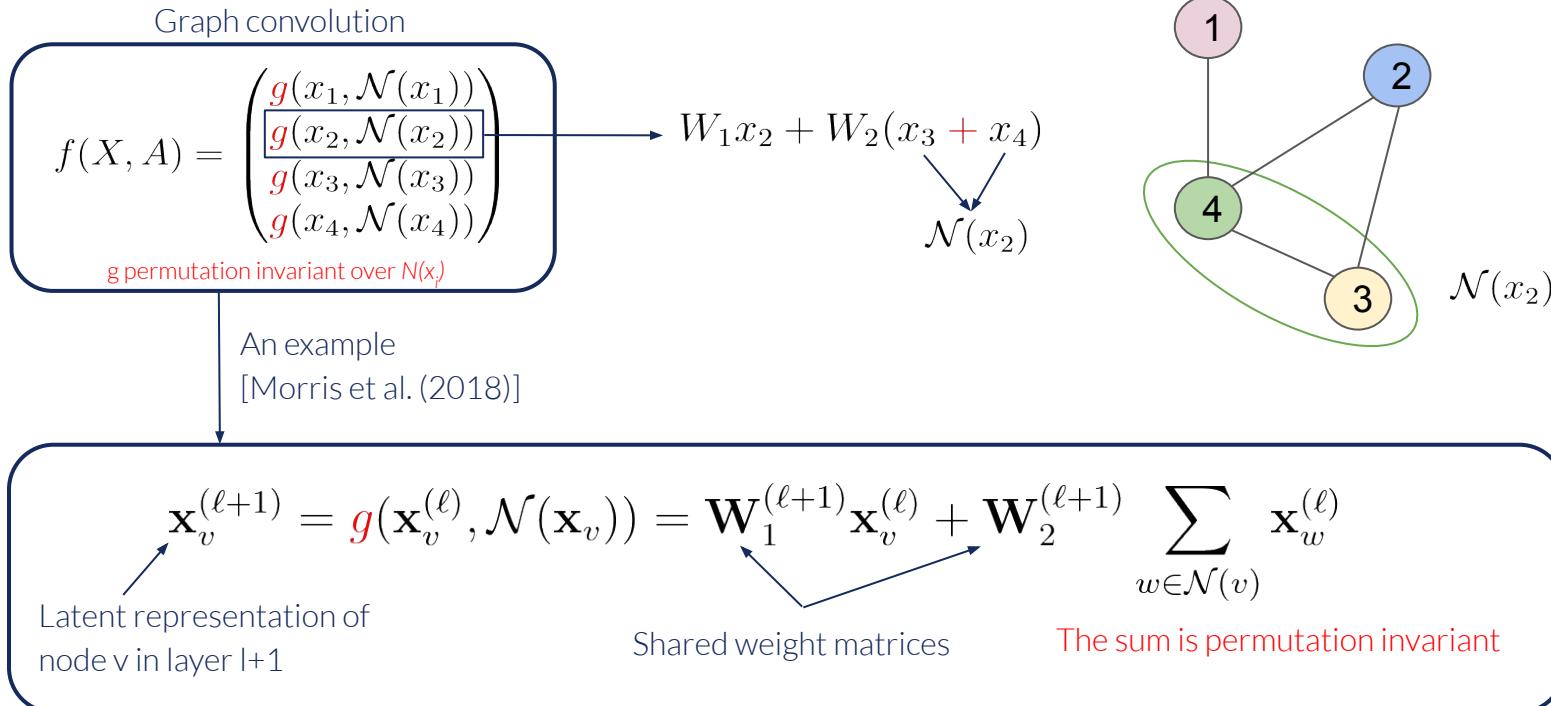
$$f(X, A) = \begin{pmatrix} g(x_1, \mathcal{N}(x_1)) \\ g(x_2, \mathcal{N}(x_2)) \\ g(x_3, \mathcal{N}(x_3)) \\ g(x_4, \mathcal{N}(x_4)) \end{pmatrix} \rightarrow W_1 x_2 + W_2 (x_3 + x_4)$$

g permutation invariant over $\mathcal{N}(x_i)$

Network parameters

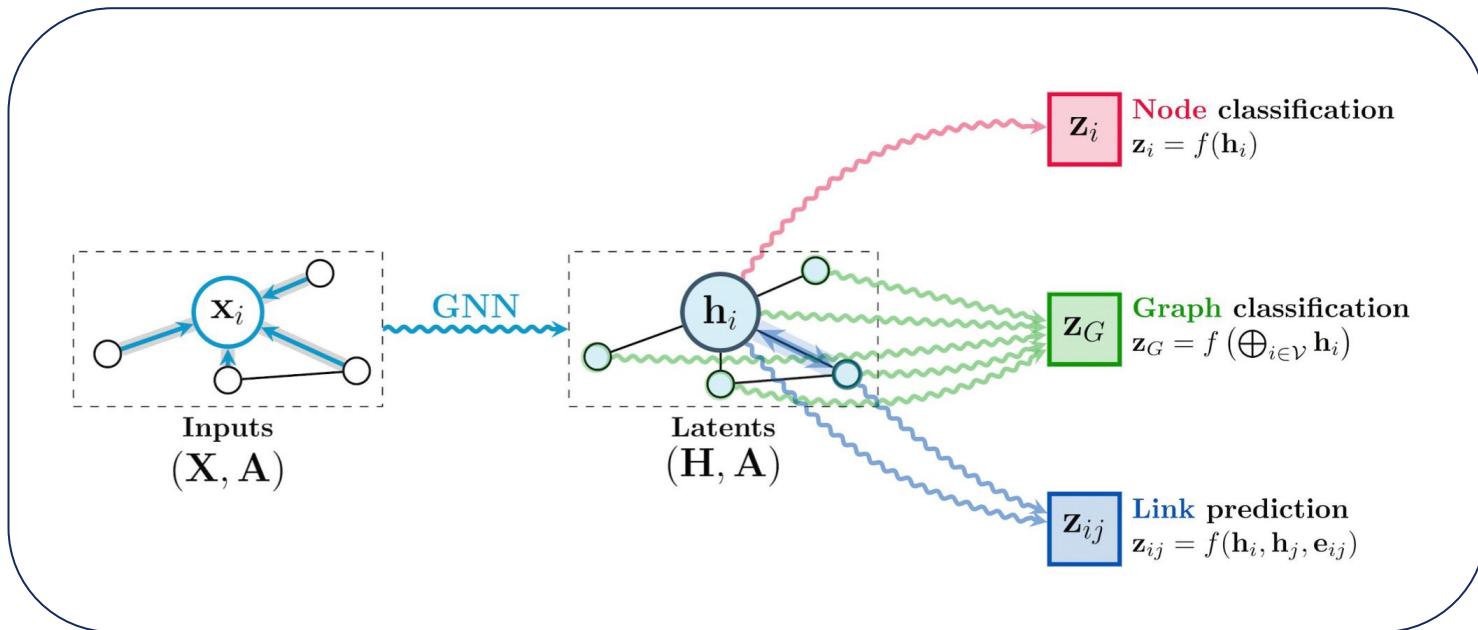


The inner workings of Graph Neural Networks



Graph Neural Networks solve three different tasks

Depending on the specific learning task, the latent nodes representations are aggregated differently

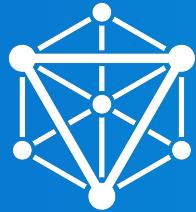


Let's get our hands dirty with some GNNs!



Colab Notebook

1. Implement a **Graph-convolution** module
2. Build a deep network for **graph classification**
3. **Train** the model
4. **Evaluate** the model



Interpretability

How to evaluate the outcome of your model

The difference between inherent and post hoc interpretability

 **Interpretability** is the degree to which a human can understand the cause of a decision

- Some models are **inherently explainable**, which means that by definition we have a complete understanding of their mechanics
 - Linear and logistic regression
 - Decision trees
 - Nearest neighbors
- Increasingly complex models, as *Deep Neural Networks*, are harder to grasp
 - ➔ We have to use **post-hoc interpretation**
 - ➔ Freedom for developers with **model-agnostic methods**

Current post-hoc Neural Network interpretability techniques

We can group interpretability methods for neural networks into **two main families**

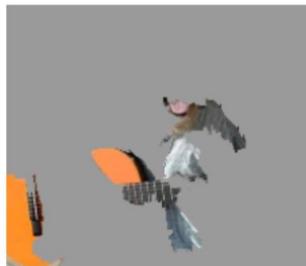
- **Proxy models** formulate *locally faithful approximations* around the prediction, using simple models (linear) or sets of rules
→ LIME (Ribeiro et al., 2016)



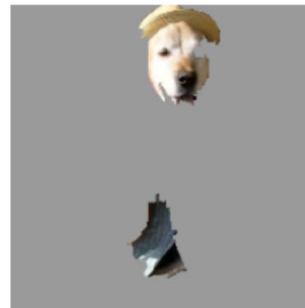
(a) Original Image



(b) Explaining *Electric guitar*



(c) Explaining *Acoustic guitar*



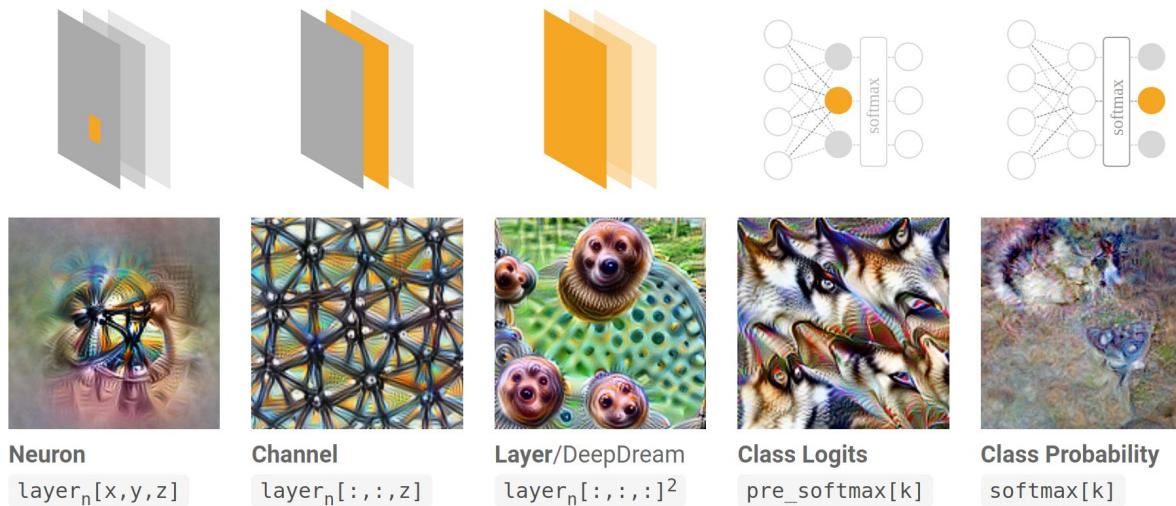
(d) Explaining *Labrador*

Fig. from "Why should i trust you?" *Explaining the predictions of any classifier*. Ribeiro et al. 2016, [arxiv:1602.04938](https://arxiv.org/abs/1602.04938)

What's going on under the hood of my ANN?

We can group interpretability methods for neural networks into **two main families**

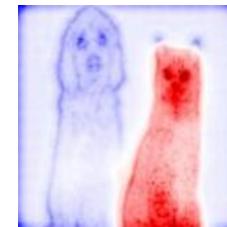
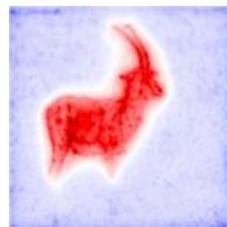
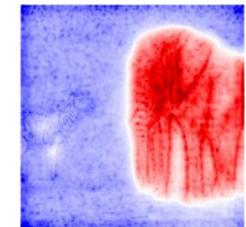
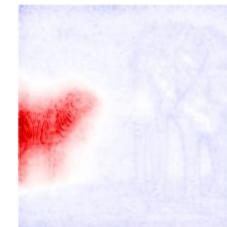
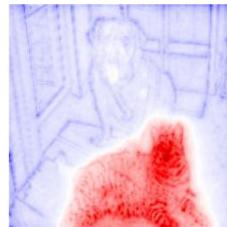
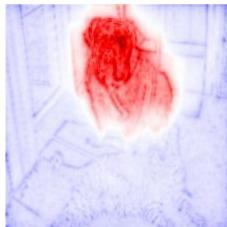
- Identify important aspects of the **computation**
 - **Learned features** through *activation maximization*
 - Find **input-specific units** through *network dissection*



What's going on under the hood of my ANN?

We can group interpretability methods for neural networks into **two main families**

- Identify important aspects of the **computation**
 - Obtain **per-class explanations**, in the form of heatmaps, using gradients and salient features



VISIUM

Graphs require different explanations for different tasks

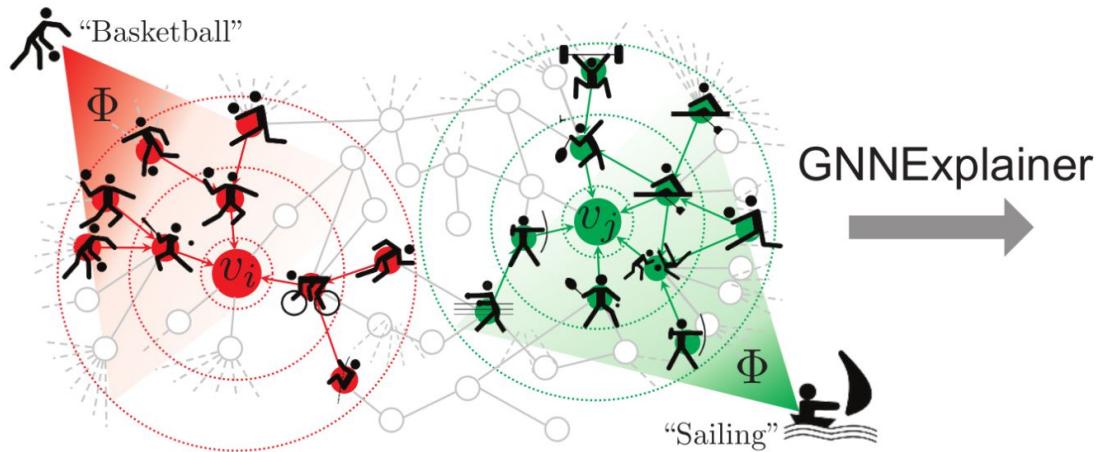
Explain a prediction from a GNN as a **small subgraph** of the input graph with a **small subset of node features**

- Node classification and Edge prediction:
 - Which **combination of neighbors** most influence the current one?
E.g. Origin of traffic jam on a specific road (edge) or intersection (node)
- Graph classification:
 - Which **part of the graph**, and which **properties**, determine the output?
E.g. Which groups in a molecule characterize the physical properties

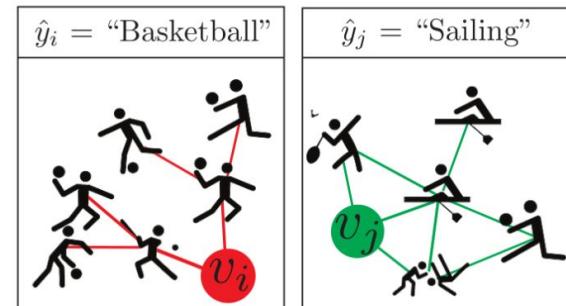
How to handle interpretability on graphs? → GNN Explainer

Explain a prediction from a GNN as a **small subgraph** of the input graph with a **small subset** of node features

GNN model training and predictions



Explaining GNN's predictions



Optimizing mutual information to explain predictions

The explainer objective is to **maximize** the **mutual information (MI)** between the network prediction Y over the full computation graph G_C and the prediction conditioned over a subgraph G_S and subsets of features X_S^F

$$\max_{G_S, X_S^F} MI(Y, (G_S, X_S^F)) = H(Y) - H(y|G = G_S, X = X_S^F)$$

We can reformulate as a **cross entropy minimization** for a graph mask M and a feature mask F

$$\min_{M, F} - \sum_{c=1}^C \mathbf{1}_{y=c} \log P_\Phi(Y = y | G = A_c \odot \sigma(M), X = X_c \odot F)$$

And include **penalization terms**:

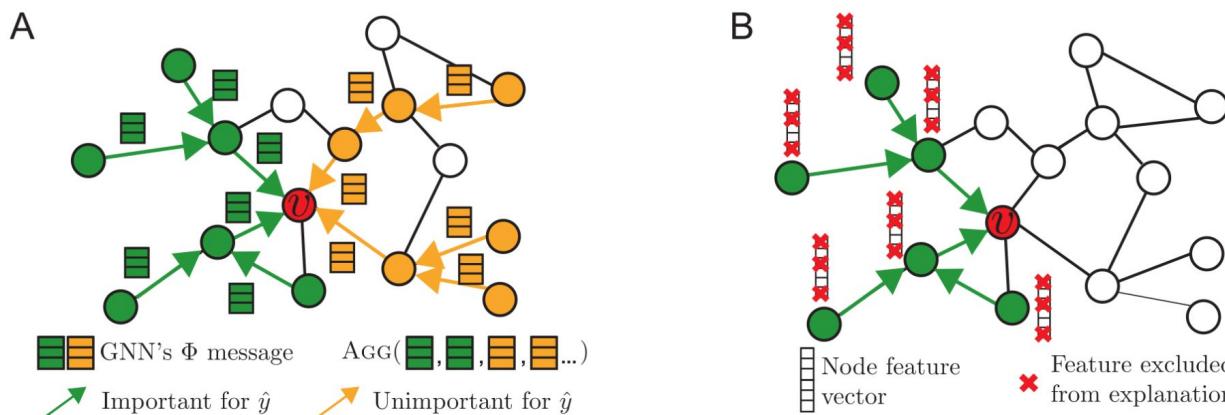
- Constraint on the **graph size** $\lambda_1 \|\sigma(M)\|_1$
- Constraint on the **graph entropy** $\lambda_2 H(\sigma(M))$
- Constraint on the **number of features** $\lambda_3 \|\sigma(F)\|_1$

Let's get our hands dirty with GNN Explainer!

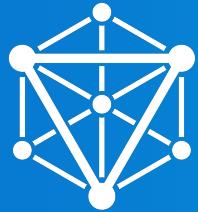


Colab Notebook

1. Implement the **Loss function**
2. Understand **GNN Explainer** workflow
3. **Explain** some graphs



VISIUM



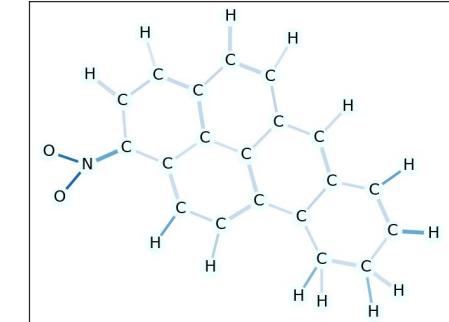
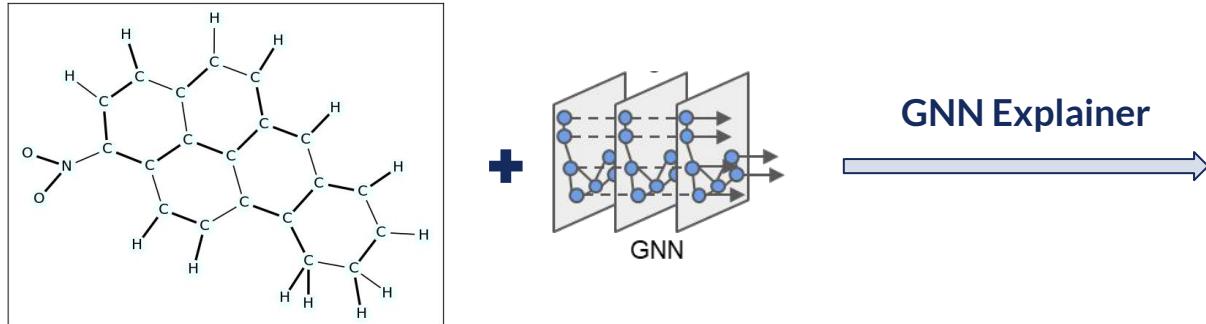
Summary

What we learned today

What you learned today

During the course of this workshop, we learned about:

- **Graph data** and how to represent them;
- How to do **deep learning** on graphs;
- The importance of **interpretability** in ML, and
- How to explain a GNN prediction using **GNN Explainer**.



Questions?



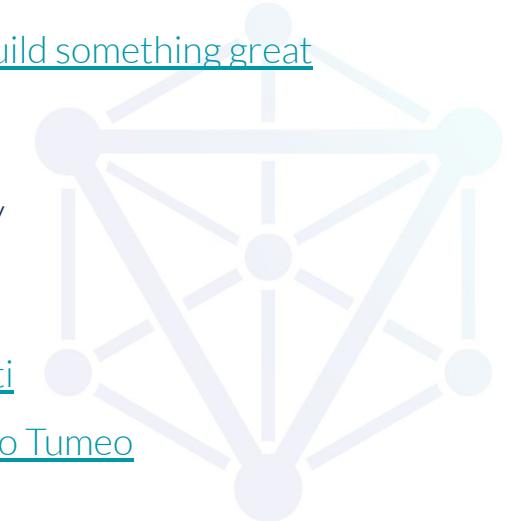
Thank you for your participation and attendance!

For more information about **Visium**

- [Visium homepage](#)
- [Work with us to build something great](#)
- [Join our team](#)

You can reach us directly

- [Arnaud Dhaene](#)
- [William Cappelletti](#)
- [Camilla Casamento Tumeo](#)



VISIUM

EPFL Innovation Park, 1015 Lausanne

Technopark Zurich, 8005 Zurich



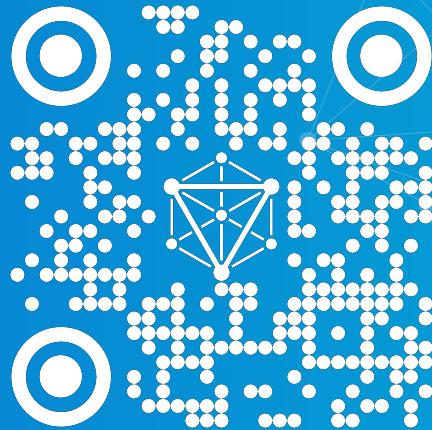
www.visium.ch



+41 21 691 55 30



arnaud.dhaene@visium.ch



VISIUM